

The Compression Paradox in AI

Why Meaning Can Degrade Before Models Hallucinate

SFP 03 | Semantic Fidelity Project

A. Jacobs

Abstract

Generative AI depends on compression. It reduces large bodies of information into representations that can be retrieved, summarized, and regenerated at scale. The same process can remove context, hierarchy, causal structure, and provenance while leaving the output fluent and recognizable.

This paper defines the Compression Paradox as the condition in which transformations that appear safe because they reduce visible error can produce less visible losses of meaning. It distinguishes compression from drift and explains how generation can conceal what compression removed.

Core Claim

Compression creates semantic risk when systems preserve recognizable content while removing the relationships that made the content meaningful.

Compression as a Necessary Operation

Intelligence requires compression. Human memory, language, institutional records, and machine learning systems all reduce complexity into reusable forms. Without compression, information cannot be stored, transmitted, compared, or applied at scale.

The problem is therefore not compression itself. A good representation removes detail while preserving the relationships relevant to its purpose. A map omits most features of a territory and remains useful because it preserves the spatial relationships needed for navigation. Semantic loss begins when the system cannot distinguish expendable detail from meaning-bearing structure.

The Compression Paradox

The Compression Paradox describes a recurring property of generative systems. Tasks such as summarization, simplification, synthesis, and paraphrase appear safe because they rarely require the model to invent new facts. Yet these tasks can alter meaning before any recognizable hallucination occurs.

The output may retain names, claims, and topical similarity while losing causal order, uncertainty, exception, emphasis, or provenance. Because the result remains fluent, the loss is interpreted as

successful simplification. The transformation looks safer precisely because its errors are structural rather than spectacular.

What Compression Removes First

Early compression usually removes redundancy and surface detail. Continued compression begins to remove context, weak signals, competing interpretations, and the relative weight of claims. Later transformations may preserve conclusions while severing them from the evidence and conditions that justified them.

Different domains lose different structures. Scientific material loses uncertainty and methodological limits. Legal material loses exceptions and scope. Institutional material loses the sequence by which decisions were made. Personal communication loses tone, history, and implied purpose. In every case, the representation can remain readable after the structure needed for faithful interpretation has thinned.

Recursive Compression

Recursive compression occurs when later representations are built from earlier ones rather than returning to the original source. A summary may be condensed again, or a generated explanation may eventually become part of the material used to produce something new. Over time, each stage begins from information that has already been selected, simplified, and partially discarded.

Recursive compression can still support learning when systems retain access to sources and reality can correct their models. It becomes a driver of semantic drift when derivative representations circulate more easily than the original referent. Later stages then optimize inherited simplifications rather than encounter the conditions from which the information arose.

Generation Conceals Deletion

Traditional lossy compression leaves visible traces. A low-resolution image looks degraded. Generative reconstruction behaves differently. It fills gaps with plausible structure, smooths transitions, and produces a complete surface. The reader sees a finished account rather than the absence on which it was built.

This creates the most important asymmetry in the Compression Paradox. Compression removes structure. Generation conceals that it was removed. Fluency is therefore weak evidence that the underlying relationships survived.

Why Benchmarks Miss the Loss

Benchmarks usually reward concise coverage, factual support, similarity to a reference answer, or human preference. These measures can favor outputs that reduce ambiguity and make conclusions easier to consume. The very features that improve the score may lower fidelity when ambiguity or qualification belongs to the source.

A fidelity test must examine what disappeared and whether the loss changes interpretation. It should compare the transformation with the original source, identify missing relationships, and test whether those relationships can still be recovered. Measuring only the finished output cannot reveal how much of its apparent completeness was reconstructed.

Preserving Fidelity Under Compression

Systems can compress without treating every omitted element as irrelevant. They can preserve links to sources, distinguish quotation from reconstruction, carry uncertainty forward, record exclusions that materially affect interpretation, and allow users to expand compressed claims back into their supporting context.

The correct goal is not maximal preservation. No representation can carry everything. The goal is controlled loss, with the preserved structure determined by the purpose of the representation and with enough provenance to restore contact when the task changes.

Conclusion

Hallucination is visible because the system adds something that is not supported. Compression failure is harder to see because the system removes something and then generates a surface that no longer reveals the absence. Semantic fidelity provides a way to examine that earlier failure point.