

Why ChatGPT Can Sound Right While Getting the Answer Wrong

Fluent Reconstruction and the Appearance of Understanding

A. Jacobs | Semantic Fidelity Project

Series Context

This paper applies the definitions in *What Is Semantic Fidelity*, *When Accuracy Is Not Enough*, and *The Compression Paradox in AI*. Semantic fidelity concerns the preservation of meaning across transformation. Semantic drift is gradual change in that meaning. Reality Drift is a broader system condition in which representations continue organizing action while losing corrective contact with what they represent.

Shared Illustrative Case

The following report is fictional and isolates the mechanism. A source report says:

“In this small observational sample, Clinic A had fewer repeat visits after introducing remote follow-up. Patients also received additional staffing support. The report cannot determine which change produced the difference, and the result should not be generalized to other clinics without further study.”

A later summary says:

“Remote follow-up reduced repeat visits at Clinic A.”

This paper examines one point at which the later sentence can be produced, accepted, or reused.

Core Claim

A well-formed answer can conceal which parts came from the source, which were inferred, and which were supplied by a familiar pattern.

An assistant receives the clinic report and a user’s question: *“Should our clinic adopt remote follow-up to reduce repeat visits?”* It replies: *“Yes. The report shows that remote follow-up reduced repeat visits at Clinic A, so adopting it is a reasonable evidence-based step.”* The response is direct. It sounds responsive to the user’s goal. It also changes a limited observation into a causal finding and then into a recommendation for a different clinic.

The answer crosses three boundaries. It moves from source to interpretation when it treats a decline after two simultaneous changes as evidence of one intervention’s effect. It moves from interpretation to generalization when it assumes Clinic A’s result applies elsewhere. It moves from

generalization to advice when it recommends adoption. Fluency can hide these crossings because the answer presents them as one continuous thought.

Language models can draw on source material, conversation context, learned patterns, and tool results. The important question here is not whether every response can be explained by one mechanism such as token prediction. It is whether the particular answer kept its inferences constrained by the report and by the user's actual decision. The source should be able to resist the answer. In this case its uncertainty about causation and its warning against generalization should interrupt the recommendation.

There are several ways the same surface impression can arise. The model may miss the staffing sentence. It may register the sentence but give greater weight to the apparent intervention-result pair. It may answer the user's implied request for a decision rather than explain that the report cannot resolve it. It may produce a useful-sounding recommendation from general expectations about remote follow-up. These are different diagnoses. Their common feature is a clean output whose apparent completeness exceeds what its evidence supports.

The phrase synthetic coherence is useful here when an answer's smooth structure and internal consistency make it appear more complete than its grounding warrants. It does not mean that every fluent answer is unreliable. It identifies an evidentiary gap that style can conceal. A reader should be able to separate what the report observed, what the assistant inferred, and what additional information a decision would require.

A better answer might say: *"The report describes fewer repeat visits after remote follow-up and additional staffing were introduced at one clinic. It cannot show which change caused the decline or whether the result would transfer to your clinic. It supports considering a pilot and deciding in advance what outcome and comparison would test it."* This remains useful. The report's limits shape the recommendation instead of disappearing beneath it.

The same failure can spread through tools and workflows. An assistant's answer becomes a meeting brief. A second assistant uses that brief as context. A later plan states that the intervention has already been shown to work. The first answer need not contain a dramatic hallucination to alter the evidence available to every later step. Its unsupported transition becomes part of the next system's input.

Diagnostic Question

Which statements are observations from the source, which are inferences, and which new decision does the answer make on the user's behalf?

Connection to Reality Drift

A wrong answer is not by itself the full system condition. Reality Drift becomes a concern when coherent answers are repeatedly reused as evidence while the source's ability to correct them weakens. Coherence is not contact.