

Why Retrieval and Citations Can Still Lose the Meaning of a Source

Access, Selection, Interpretation, and Corrective Force in RAG

A. Jacobs | Semantic Fidelity Project

Series Context

This paper applies the definitions in *What Is Semantic Fidelity*, *When Accuracy Is Not Enough*, and *The Compression Paradox in AI*. Semantic fidelity concerns the preservation of meaning across transformation. Semantic drift is gradual change in that meaning. Reality Drift is a broader system condition in which representations continue organizing action while losing corrective contact with what they represent.

Shared Illustrative Case

The following report is fictional and isolates the mechanism. A source report says:

“In this small observational sample, Clinic A had fewer repeat visits after introducing remote follow-up. Patients also received additional staffing support. The report cannot determine which change produced the difference, and the result should not be generalized to other clinics without further study.”

A later summary says:

“Remote follow-up reduced repeat visits at Clinic A.”

This paper examines one point at which the later sentence can be produced, accepted, or reused.

Core Claim

Retrieval gives an answer access to a source. Semantic fidelity depends on whether the source’s decisive relationships actually constrain the answer.

Imagine the clinic report in a searchable repository. A retrieval-augmented system is asked whether remote follow-up reduces repeat visits. The report is returned among the top results. The final answer states, *“Yes. A clinic report found that remote follow-up reduced repeat visits,”* and cites the report. The citation is real. The document is relevant. The answer still exceeds what that document supports.

Several stages need separate inspection. Retrieval asks whether the right document was found. Selection asks whether the passages that control interpretation entered the model’s usable context. Interpretation asks whether the answer preserved the source’s scope, uncertainty, and causal limits. Attribution asks whether the cited passage actually supports the attached claim. Reuse asks what

happens when the generated answer is treated as a new source. A system can succeed at retrieval and fail at any later stage.

Similarity can help find the clinic report because it shares language with the query. It does not settle what the report means. A passage about remote follow-up and repeat visits may rank highly while a later passage explains that staffing also changed. If only the first passage is selected, the system may never encounter the qualification. If both enter context, the model can still produce the causal answer. These are different failures with different remedies. More retrieval does little for an interpretation problem. Better interpretation cannot repair a decisive passage that never entered context.

The citation introduces another possibility. A reader sees a linked report and treats it as evidence that the claim was checked. Yet a citation records a proposed relationship between an answer and a source. It does not establish that relationship. In this case the cited source should correct “*remote follow-up reduced repeat visits*” to a narrower claim. If neither the system nor its readers permit that correction, the citation can increase confidence in the wrong inference.

This is the difference between source presence and source force. Presence means the report appears in the workflow. Force means its contrary or limiting details can change the generated answer. A source that is retrieved, displayed, and cited but cannot change the conclusion has become evidence in appearance while losing some of its corrective role.

The relevant evaluation can follow the claim backward. Locate the exact sentence supporting “*reduced*.” Check whether it establishes causation, whether another passage limits it, and whether the final answer carries those limits. Then follow the answer forward. Does a later summary cite the model’s phrasing or the report itself? Can a reader recover the full source? Does correction to the report’s interpretation propagate to the derivative answer? This is where provenance and semantic fidelity meet. Knowing where an answer came from helps, but the lineage also needs to preserve what the source permitted the answer to say.

Diagnostic Question

If the cited source disagrees with the answer’s scope or certainty, where in the workflow can it change the answer?

Connection to Reality Drift

Correct retrieval alongside unsupported interpretation is a local failure. Reality Drift becomes a concern when systems repeatedly produce sourced-looking answers while the underlying sources lose the ability to constrain or correct those answers. The system can remain operational and apparently grounded even as corrective contact weakens.